



Reti di Calcolatori e Cybersecurity

# Routing Distance Vector

Ing. Vincenzo Abate

# Algoritmo Distance Vector

Ogni router mantiene una tabella di tutti gli instradamenti noti

- inizialmente, solo le reti a cui è connesso direttamente

Ogni entry della tabella indica:

- una rete raggiungibile
  - il next hop
  - il numero di hop necessari per raggiungere la destinazione
- Ch' raggiunge, con quanti passi e chi è il prossimo?*

Periodicamente, ogni router invia a tutti i vicini (due router sono vicini se sono collegati alla stessa rete fisica): *info su tutte le dest. e le mando ai miei vicini*

- un messaggio di aggiornamento contenente tutte le informazioni della propria tabella (vettore delle distanze – **distance vector**)

I router che ricevono tale messaggio aggiornano la tabella nel seguente modo:

- eventuale modifica di informazioni relative a cammini già noti
- eventuale aggiunta di nuovi cammini
- eventuale eliminazione di cammini non più disponibili

# Algoritmo Distance Vector

Destin.	Dist.	Route
net 1	0	direct
net 2	0	direct
net 4	8	router L
net 17	5	router M
net 24	6	router A
net 30	2	router Q
net 42	2	router A

Tabella del router B

Gasto come hop.

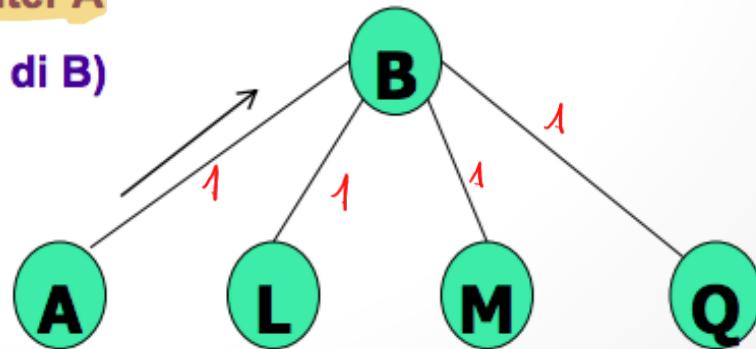
Destin.	Dist.
net 1	2
net 4	3
net 17	6
net 21	4
net 24	5
net 30	10
net 42	3

Messaggio di aggiornamento del router A

(vicino di B)

Destin.	Dist.	Route
net 1	0	direct
net 2	0	direct
net 4	4	router A
net 17	5	router M
net 24	6	router A
net 30	2	router Q
net 42	4	router A
net 21	5	router A

Tabella aggiornata del router B



# Algoritmo Distance Vector

Il calcolo **consiste nella fusione di tutti i distance vector delle linee attive**

Un **router ricalcola le sue tabelle** se:

- **cade una linea attiva**
- **riceve un distance vector, da un nodo adiacente, diverso da quello memorizzato**

Se le **tabelle risultano diverse da quelle precedenti**:

- **invia ai nodi adiacenti un nuovo distance vector**

# Algoritmo Distance Vector

Based on *Bellman-Ford* (BF) equation (dynamic programming):

Bellman-Ford equation

Let  $D_x(y)$ : cost of least-cost path from  $x$  to  $y$ .

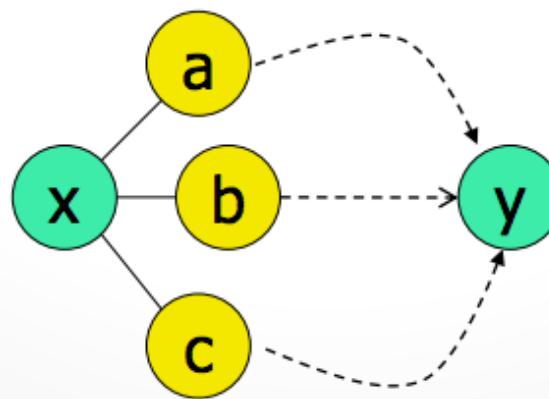
Then:

$$D_x(y) = \min_v \{ c_{x,v} + D_v(y) \}$$

$\min$  taken over all neighbors  $v$  of  $x$

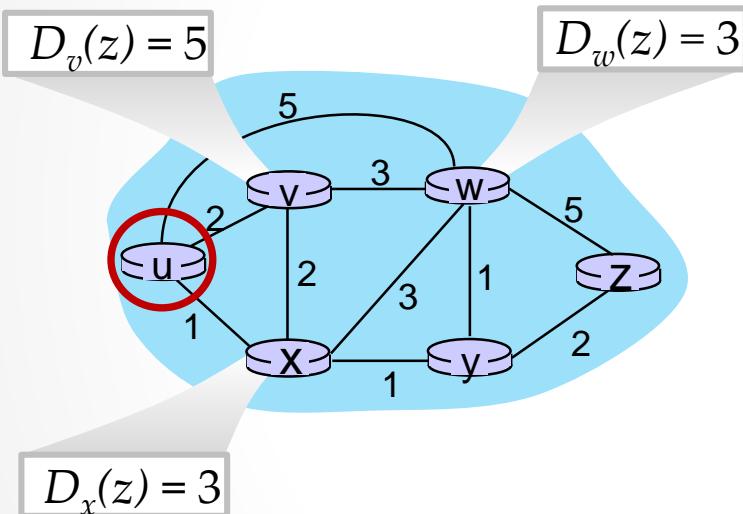
$v$ 's estimated least-cost-path to  $y$

direct cost of link from  $x$  to  $v$



# Bellman-Ford Example

Suppose that  $u$ 's neighboring nodes,  $x, v, w$ , know that for destination  $z$ :



Bellman-Ford equation says:

$$\begin{aligned} D_u(z) &= \min \{ c_{u,v} + D_v(z), \\ &\quad c_{u,x} + D_x(z), \\ &\quad c_{u,w} + D_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

node achieving minimum ( $x$ ) is next hop on estimated least-cost path to destination ( $z$ )

# Distance vector algorithm

key idea:

- from time-to-time, each node sends its own distance vector estimate to neighbors
- when  $x$  receives new DV estimate from any neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{c_{x,v} + D_v(y)\} \text{ for each node } y \in N$$

- under minor, natural conditions, the estimate  $D_x(y)$  converge to the actual least cost  $d_x(y)$

# Distance vector algorithm

## Define

- $D_x(y)$  = estimate of least cost from  $x$  to  $y$
- Distance vector:  $D_x = [D_x(y): y \in N]$

Node  $x$  knows cost to each neighbor  $v$ :  $c(x,v)$

Node  $x$  maintains  $D_x = [D_x(y): y \in N]$

Node  $x$  also maintains its neighbors' distance vectors

- For each neighbor  $v$ ,  $x$  maintains  $D_v = [D_v(y): y \in N]$

# Distance vector algorithm:

each node:

*wait* for (change in local link cost  
or msg from neighbor)

*recompute* DV estimates using DV  
received from neighbor

if DV to any destination has  
changed, *notify* neighbors

**iterative, asynchronous:** each local iteration caused by:

- local link cost change
- DV update message from neighbor

**distributed, self-stopping:** each node notifies neighbors *only* when its DV changes

- neighbors then notify their neighbors – *only if necessary*
- no notification received, no actions taken!

# Algoritmo Distance Vector

Ad ogni nodo,  $x$ :

- 1 Inizializzazione:
- 2 per tutti i nodi adiacenti  $v$ :  
     $D^x(*, v) = \text{infinito}$  {il simbolo \* significa "per ogni riga"}
- 3  $D^x(v, v) = c(x, v)$
- 4 per tutte le destinazioni,  $y$
- 5 manda  $\min_w D(y, w)$  a ogni vicino  
    ↑ ogni nodo adiacente

# Algoritmo Distance Vector

```
→ 8 loop
  9 aspetta (fino a quando vedo una modifica nel costo di un
  10 collegamento oppure ricevo un messaggio da un vicino v)
  11
  12 if (c(x,v) cambia di d)
  13 { cambia il costo a tutte le dest. via vicino v di d }
  14 { nota: d puo' essere positivo o negativo }
  15 per tutte le destinazioni y:  $D^X(y,v) = D^X(y,v) + d$ 
  16
  17 else if (ricevo mess. aggiornamento da v verso destinazione y)
  18 { cammino minimo da v a y e' cambiato }
  19 { V ha mandato un nuovo valore per il suo  $\min_w D^V(y,w)$  }
  20 { chiama questo valore "newval" }
  21 per la singola destinazione y:  $D^X(y,v) = c(x,v) + newval$ 
  22
  23 if hai un nuovo  $\min_w D^X(y,w)$  per una qualunque destinazione y
  24 manda il nuovo valore di  $\min_w D^X(y,w)$  a tutti i vicini
  25
  26 forever
```

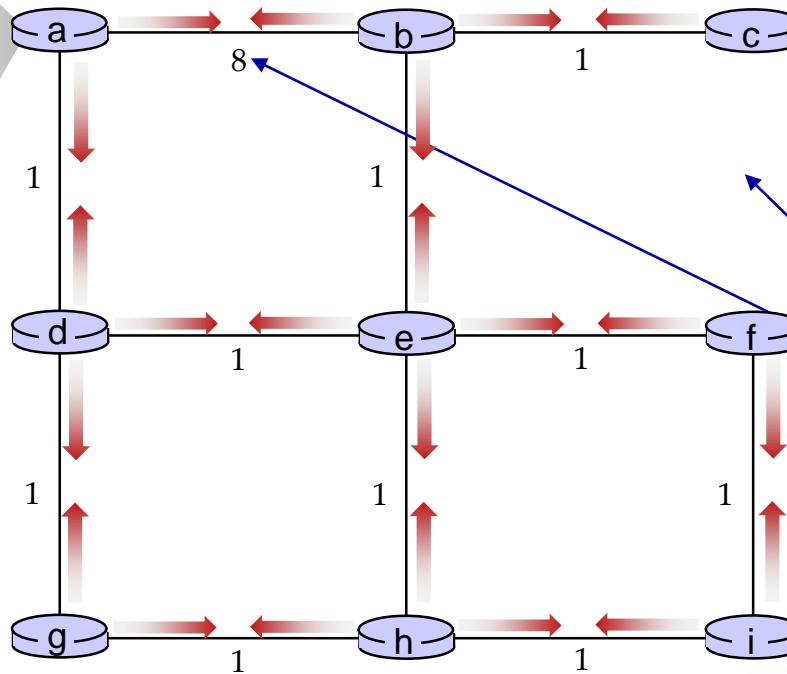
# Distance vector: example



$t=0$

- All nodes have distance estimates to nearest neighbors (only)
- All nodes send their local distance vector to their neighbors

DV in a:
$D_a(a)=0$
$D_a(b) = 8$
$D_a(c) = \infty$
$D_a(d) = 1$
$D_a(e) = \infty$
$D_a(f) = \infty$
$D_a(g) = \infty$
$D_a(h) = \infty$
$D_a(i) = \infty$



A few asymmetries:  
▪ missing link  
▪ larger cost

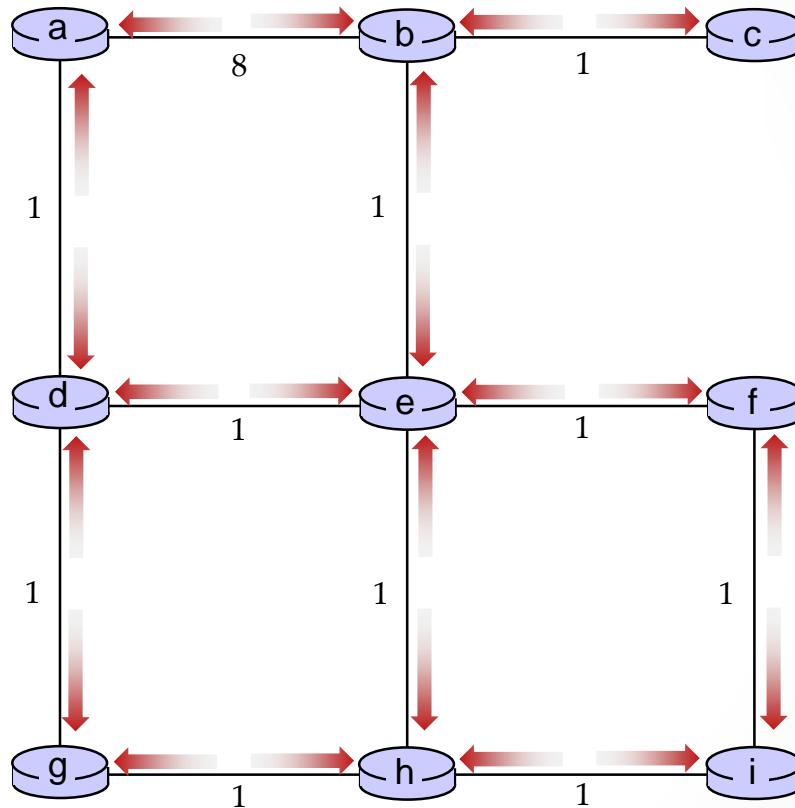
# Distance vector example: iteration



$t=1$

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors

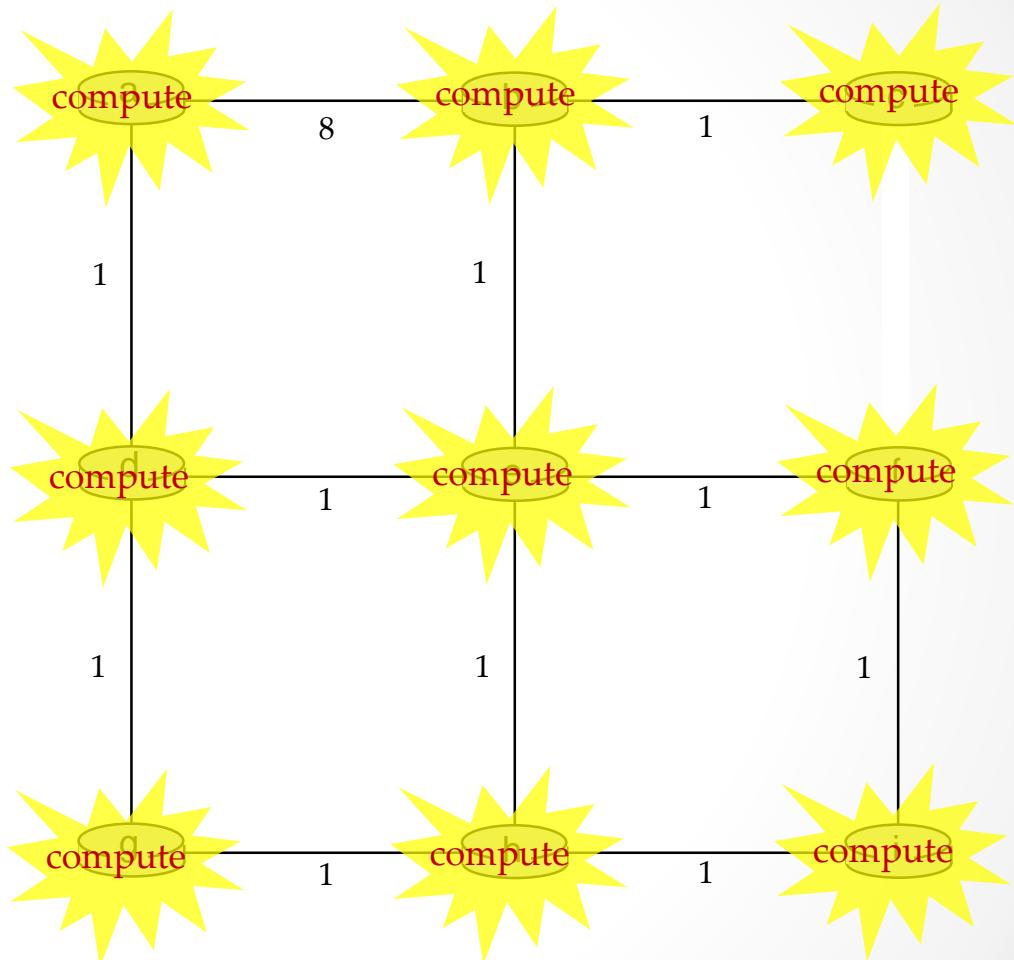


# Distance vector example: iteration



All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors



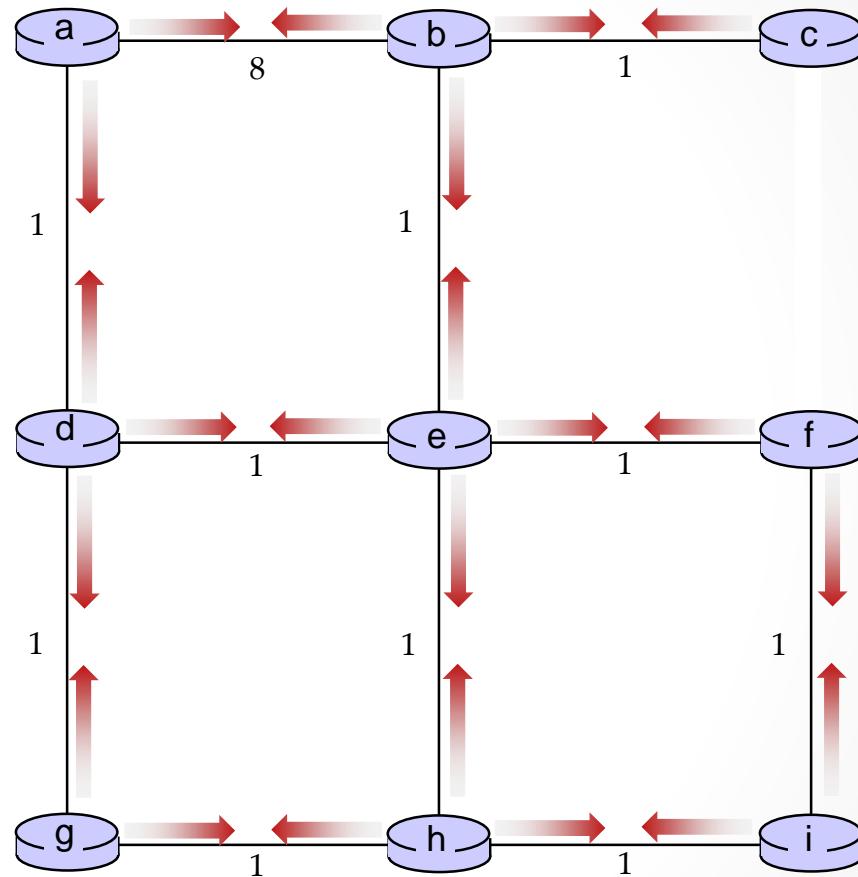
# Distance vector example: iteration



t=1

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors



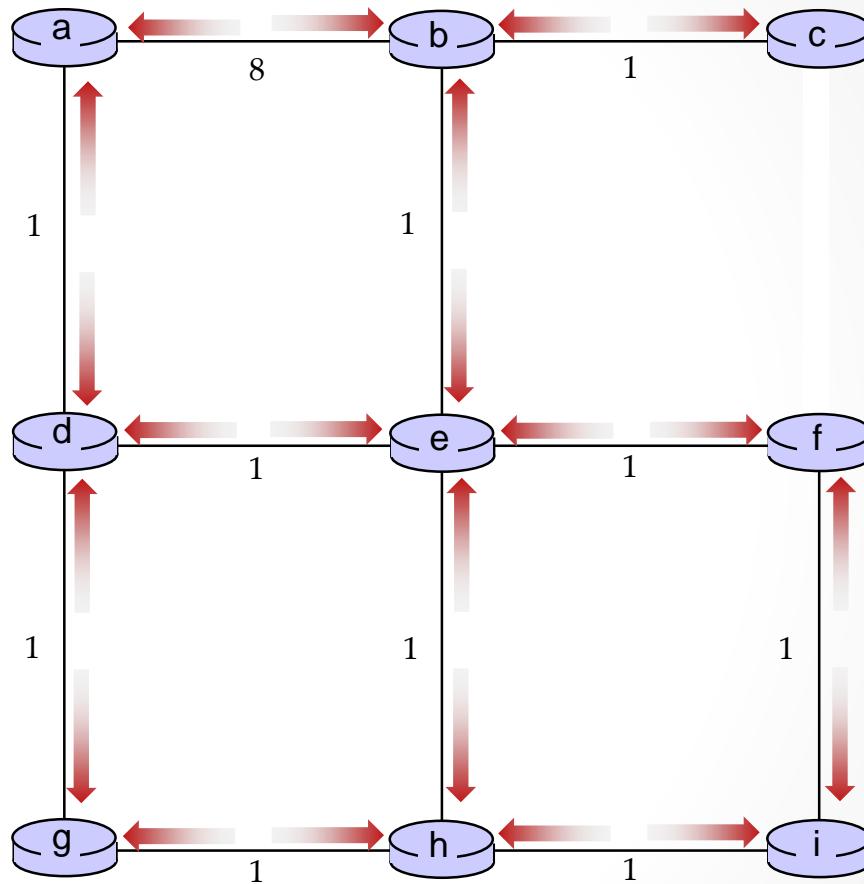
# Distance vector example: iteration



t=2

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors

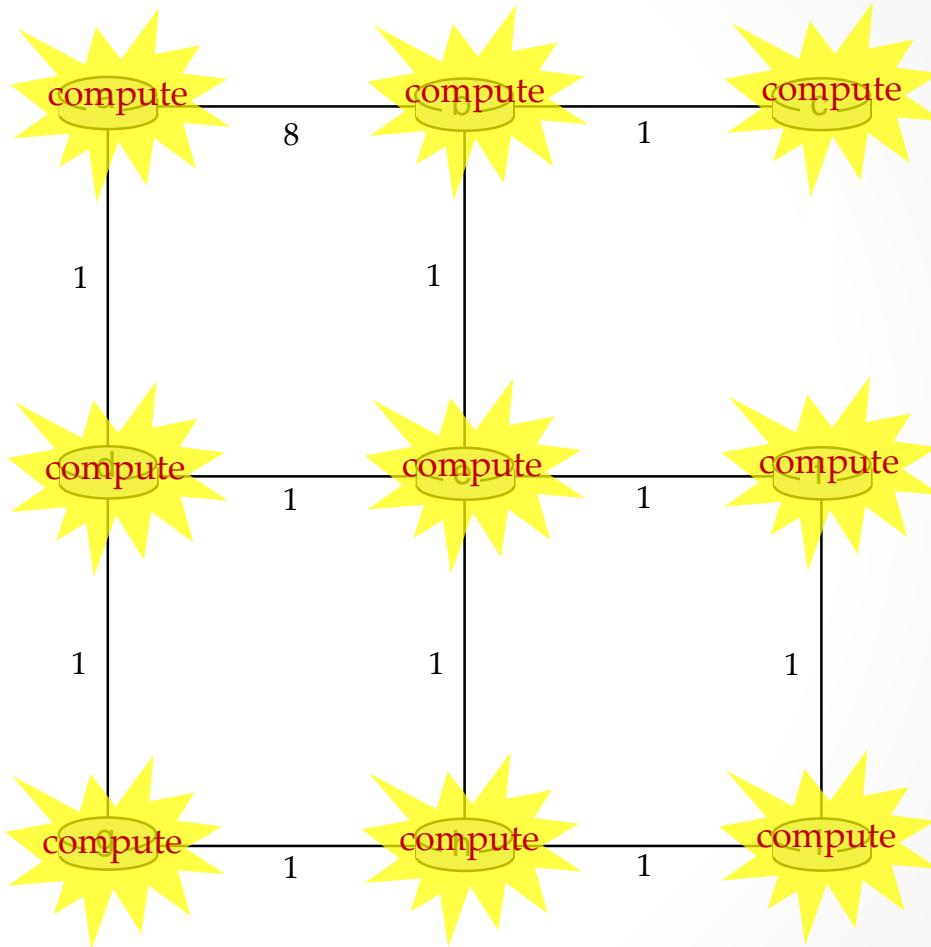


# Distance vector example: iteration



All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors



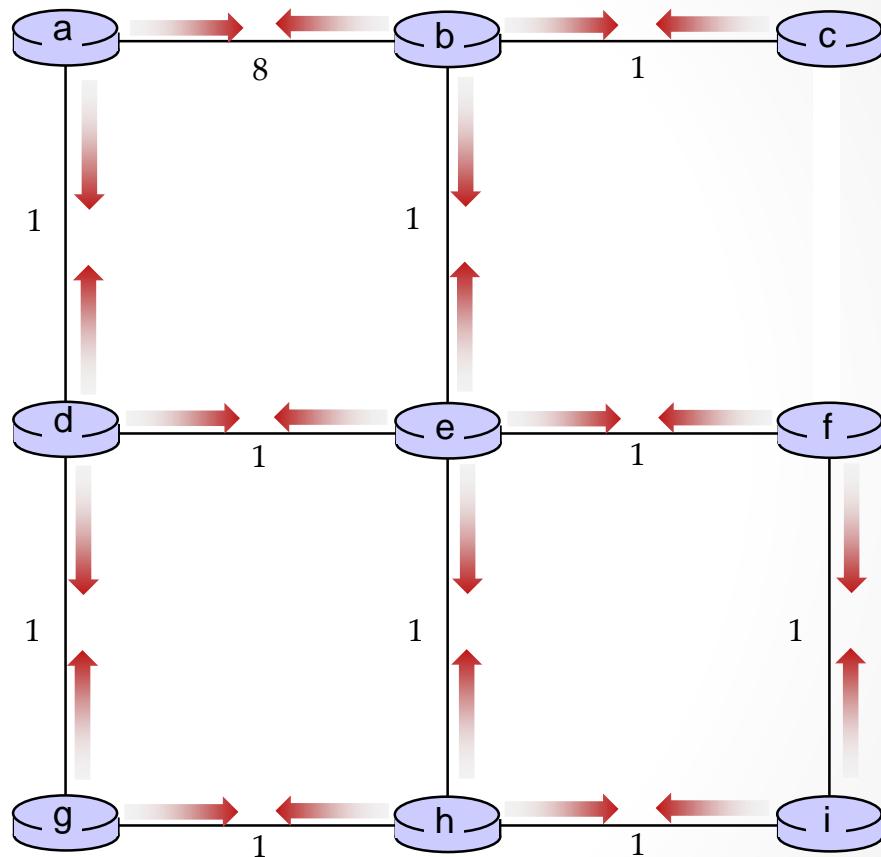
# Distance vector example: iteration



t=2

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors



# Distance vector example: iteration

.... and so on

Let's next take a look at the iterative *computations* at nodes



# Distance vector example: computation

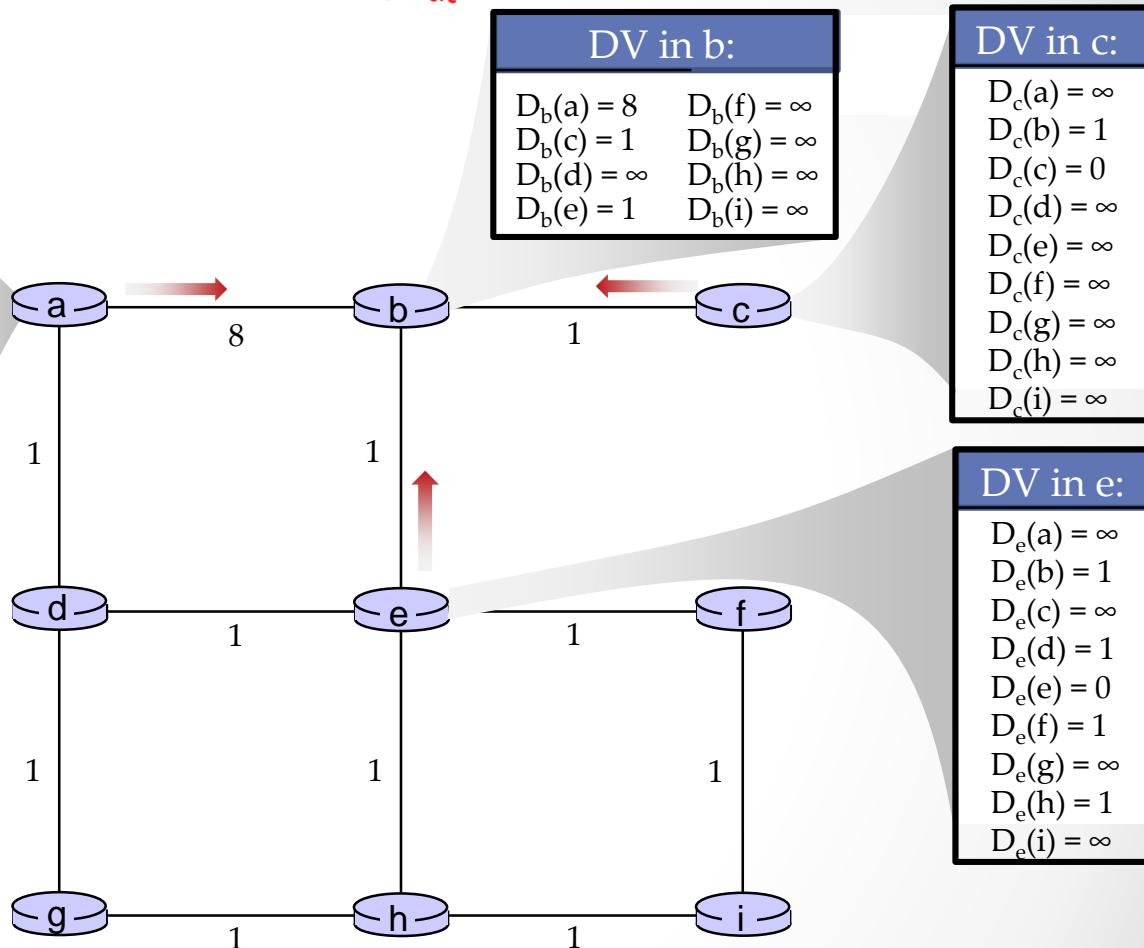


$t=1$

- b receives DVs from a, c, e

DV in a:
$D_a(a) = 0$
$D_a(b) = 8$
$D_a(c) = \infty$
$D_a(d) = 1$
$D_a(e) = \infty$
$D_a(f) = \infty$
$D_a(g) = \infty$
$D_a(h) = \infty$
$D_a(i) = \infty$

b receives dv question



# Distance vector example: computation

 t=1

- b receives DVs from a, c, e, computes:

$$\begin{aligned}
 D_b(a) &= \min\{c_{b,a} + D_a(a), c_{b,c} + D_c(a), c_{b,e} + D_e(a)\} = \min\{8, \infty, \infty\} = 8 \\
 D_b(c) &= \min\{c_{b,a} + D_a(c), c_{b,c} + D_c(c), c_{b,e} + D_e(c)\} = \min\{\infty, 1, \infty\} = 1 \\
 D_b(d) &= \min\{c_{b,a} + D_a(d), c_{b,c} + D_c(d), c_{b,e} + D_e(d)\} = \min\{9, 2, \infty\} = 2 \\
 D_b(e) &= \min\{c_{b,a} + D_a(e), c_{b,c} + D_c(e), c_{b,e} + D_e(e)\} = \min\{\infty, \infty, 1\} = 1 \\
 D_b(f) &= \min\{c_{b,a} + D_a(f), c_{b,c} + D_c(f), c_{b,e} + D_e(f)\} = \min\{\infty, \infty, 2\} = 2 \\
 D_b(g) &= \min\{c_{b,a} + D_a(g), c_{b,c} + D_c(g), c_{b,e} + D_e(g)\} = \min\{\infty, \infty, \infty\} = \infty \\
 D_b(h) &= \min\{c_{b,a} + D_a(h), c_{b,c} + D_c(h), c_{b,e} + D_e(h)\} = \min\{\infty, \infty, 2\} = 2 \\
 D_b(i) &= \min\{c_{b,a} + D_a(i), c_{b,c} + D_c(i), c_{b,e} + D_e(i)\} = \min\{\infty, \infty, \infty\} = \infty
 \end{aligned}$$

DV in a:
$D_a(a) = 0$
$D_a(b) = 8$
$D_a(c) = \infty$
$D_a(d) = 1$
$D_a(e) = \infty$
$D_a(f) = \infty$
$D_a(g) = \infty$
$D_a(h) = \infty$
$D_a(i) = \infty$

DV in b:

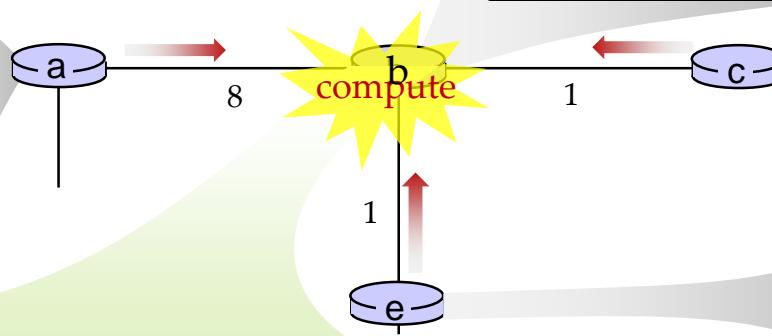
$D_b(a) = 8$	$D_b(f) = \infty$
$D_b(c) = 1$	$D_b(g) = \infty$
$D_b(d) = \infty$	$D_b(h) = \infty$
$D_b(e) = 1$	$D_b(i) = \infty$

DV in c:

$D_c(a) = \infty$
$D_c(b) = 1$
$D_c(c) = 0$
$D_c(d) = \infty$
$D_c(e) = \infty$
$D_c(f) = \infty$
$D_c(g) = \infty$
$D_c(h) = \infty$
$D_c(i) = \infty$

DV in e:

$D_e(a) = \infty$
$D_e(b) = 1$
$D_e(c) = \infty$
$D_e(d) = 1$
$D_e(e) = 0$
$D_e(f) = 1$
$D_e(g) = \infty$
$D_e(h) = 1$
$D_e(i) = \infty$



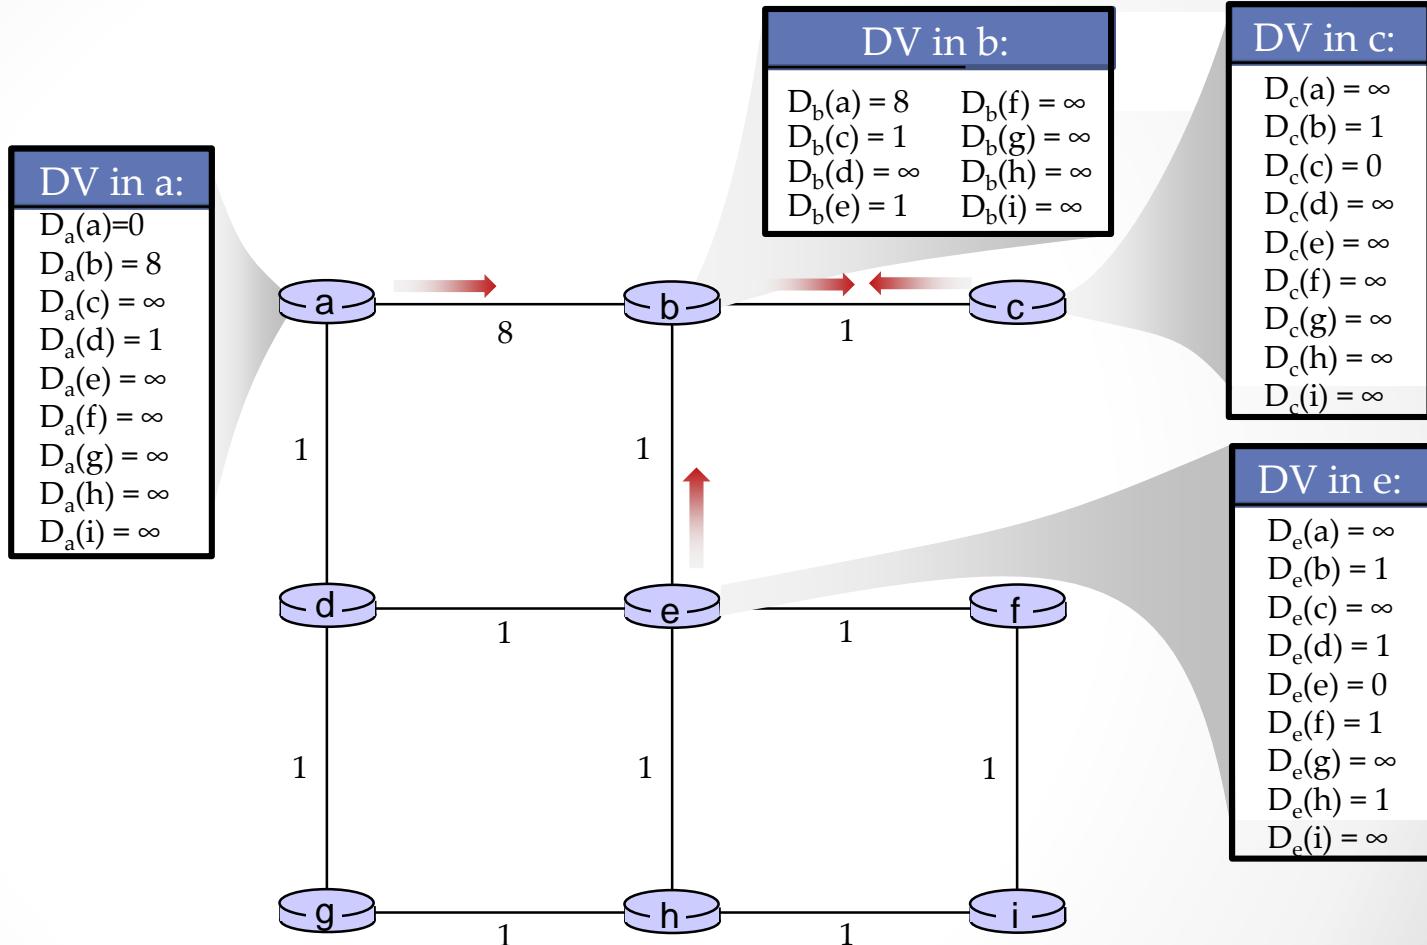
DV in b:

$D_b(a) = 8$	$D_b(f) = 2$
$D_b(c) = 1$	$D_b(g) = \infty$
$D_b(d) = 2$	$D_b(h) = 2$
$D_b(e) = 1$	$D_b(i) = \infty$

# Distance vector example: computation

 t=1

- c receives DVs from b



# Distance vector example: computation

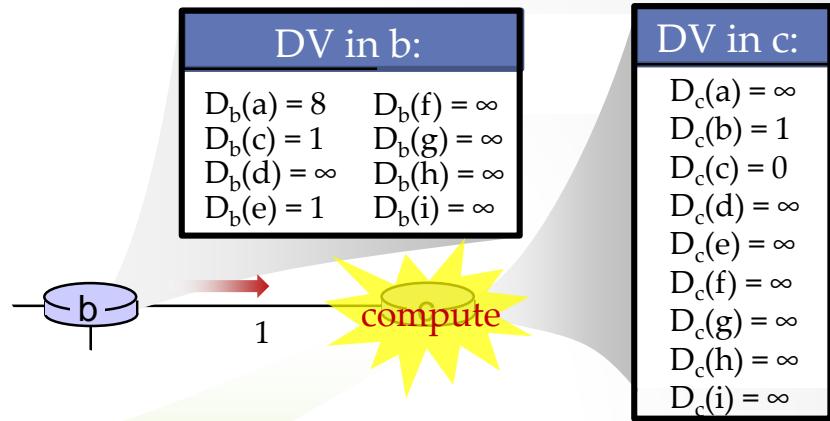
b makes prima slight adjustments



t=1

- c receives DVs from b computes:

$$\begin{aligned}D_c(a) &= \min\{c_{c,b} + D_b(a)\} = 1 + 8 = 9 \\D_c(b) &= \min\{c_{c,b} + D_b(b)\} = 1 + 0 = 1 \\D_c(d) &= \min\{c_{c,b} + D_b(d)\} = 1 + \infty = \infty \\D_c(e) &= \min\{c_{c,b} + D_b(e)\} = 1 + 1 = 2 \\D_c(f) &= \min\{c_{c,b} + D_b(f)\} = 1 + \infty = \infty \\D_c(g) &= \min\{c_{c,b} + D_b(g)\} = 1 + \infty = \infty \\D_c(h) &= \min\{c_{c,b} + D_b(h)\} = 1 + \infty = \infty \\D_c(i) &= \min\{c_{c,b} + D_b(i)\} = 1 + \infty = \infty\end{aligned}$$



DV in c:

$D_c(a) = 9$
$D_c(b) = 1$
$D_c(c) = 0$
$D_c(d) = 2$
$D_c(e) = \infty$
$D_c(f) = \infty$
$D_c(g) = \infty$
$D_c(h) = \infty$
$D_c(i) = \infty$

# Distance vector example: computation



$t=1$

- e receives DVs from b, d, f, h

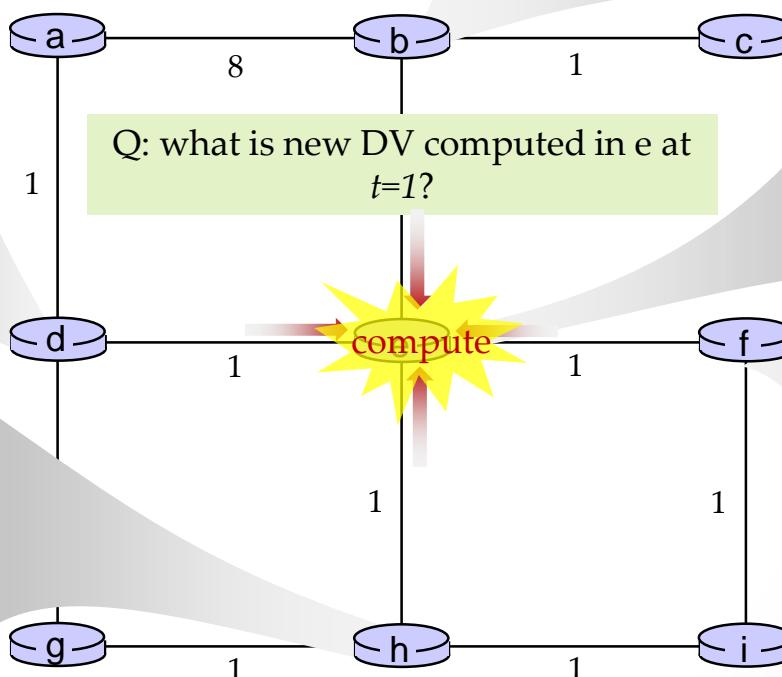
DV in d:
$D_c(a) = 1$
$D_c(b) = \infty$
$D_c(c) = \infty$
$D_c(d) = 0$
$D_c(e) = 1$
$D_c(f) = \infty$
$D_c(g) = 1$
$D_c(h) = \infty$
$D_c(i) = \infty$

DV in h:
$D_c(a) = \infty$
$D_c(b) = \infty$
$D_c(c) = \infty$
$D_c(d) = \infty$
$D_c(e) = 1$
$D_c(f) = \infty$
$D_c(g) = 1$
$D_c(h) = 0$
$D_c(i) = 1$

DV in b:
$D_b(a) = 8$
$D_b(c) = 1$
$D_b(d) = \infty$
$D_b(e) = 1$
$D_b(f) = \infty$
$D_b(g) = \infty$
$D_b(h) = \infty$
$D_b(i) = \infty$

DV in e:
$D_e(a) = \infty$
$D_e(b) = 1$
$D_e(c) = \infty$
$D_e(d) = 1$
$D_e(e) = 0$
$D_e(f) = 1$
$D_e(g) = \infty$
$D_e(h) = 1$
$D_e(i) = \infty$

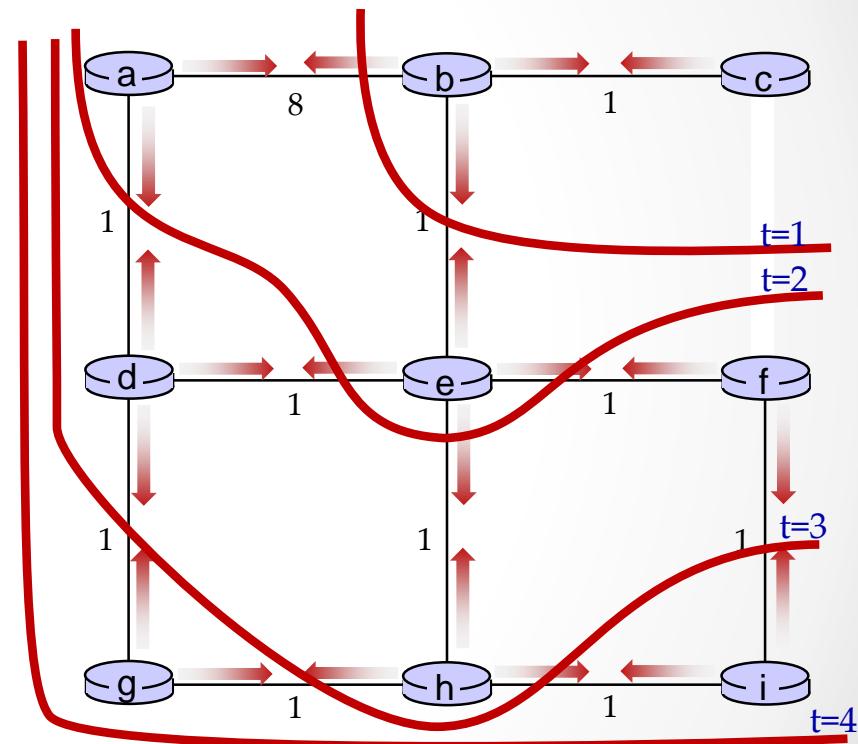
DV in f:
$D_c(a) = \infty$
$D_c(b) = \infty$
$D_c(c) = \infty$
$D_c(d) = \infty$
$D_c(e) = 1$
$D_c(f) = 0$
$D_c(g) = \infty$
$D_c(h) = \infty$
$D_c(i) = 1$



# Distance vector: state information diffusion

Iterative communication, computation steps diffuses information through network:

- ⌚ t=0 c's state at t=0 is at c only
- ⌚ t=1 c's state at t=0 has propagated to b, and may influence distance vector computations up to 1 hop away, i.e., at b
- ⌚ t=2 c's state at t=0 may now influence distance vector computations up to 2 hops away, i.e., at b and now at a, e as well
- ⌚ t=3 c's state at t=0 may influence distance vector computations up to 3 hops away, i.e., at b,a,e and now at c,f,h as well
- ⌚ t=4 c's state at t=0 may influence distance vector computations up to 4 hops away, i.e., at b,a,e, c, f, h and now at g,i as well



# Distance Vector: considerazioni

## Vantaggi:

- facile da implementare

## Svantaggi

- ogni messaggio contiene un'intera tabella di routing
- lenta propagazione delle informazioni sui cammini:
  - converge alla velocità del router più lento
- se lo stato della rete cambia velocemente, le rotte possono risultare inconsistenti
  - possono innescarsi dei loop a causa di particolari variazioni della topologia
- difficile capirne e prevederne il comportamento su reti grandi
  - nessun nodo ha una mappa della rete!

# Distance Vector: considerazioni

How fast the routers learn about link-status change  
in the network

With distance vector routing

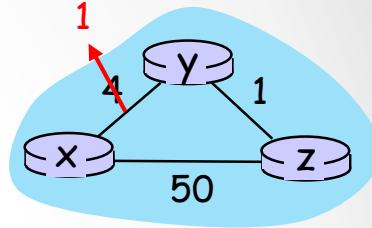
- Good news travels fast
- Bad news travels slow



# Distance vector: link cost changes

## link cost changes:

- node detects local link cost change
- updates routing info, recalculates local DV
- if DV changes, notify neighbors



$t_0$ :  $y$  detects link-cost change, updates its DV, informs its neighbors.

“good news travels fast”

$t_1$ :  $z$  receives update from  $y$ , updates its table, computes new least cost to  $x$ , sends its neighbors its DV. *Cambiaron los costos. Vaya que un menor costo*

$t_2$ :  $y$  receives  $z$ 's update, updates its distance table.  $y$ 's least costs do not change, so  $y$  does *not* send a message to  $z$ .

# Distance vector: link cost changes

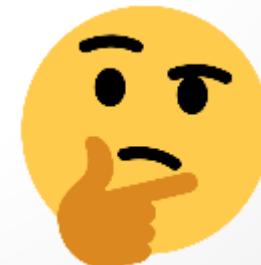
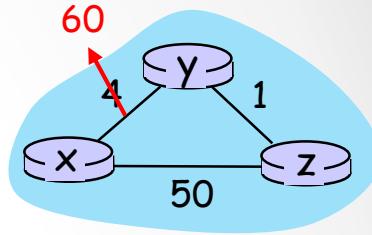
## link cost changes:

- node detects local link cost change
- “bad news travels slow” – count-to-infinity problem:

- $y$  sees direct link to  $x$  has new cost 60, but  $z$  has said it has a path at cost of 5. So  $y$  computes “my new cost to  $x$  will be 6, via  $z$ ); notifies  $z$  of new cost of 6 to  $x$ .
- $z$  learns that path to  $x$  via  $y$  has new cost 6, so  $z$  computes “my new cost to  $x$  will be 7 via  $y$ ), notifies  $y$  of new cost of 7 to  $x$ .
- $y$  learns that path to  $x$  via  $z$  has new cost 7, so  $y$  computes “my new cost to  $x$  will be 8 via  $y$ ), notifies  $z$  of new cost of 8 to  $x$ .
- $z$  learns that path to  $x$  via  $y$  has new cost 8, so  $z$  computes “my new cost to  $x$  will be 9 via  $y$ ), notifies  $y$  of new cost of 9 to  $x$ .

...

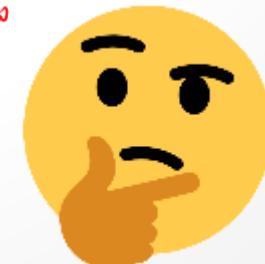
- Possible soluzione al problema?



# Count to Infinity

	1	2	3	4	Initially
A	B	C	D	E	5 Nodes = count
	1 Va ad 00	2	3	4	After 1 exchange
ora C su che B è 3	3 Si aggiorna l'ut	4 Si aggiorna l'ut	3	4	After 2 exchanges
	5	4	5	4	After 3 exchanges
	5	6	5	6	After 4 exchanges
	7	6	7	6	After 5 exchanges

Tutti quelli che arrivano ad A passano per B aggiornano  
etc... to infinity

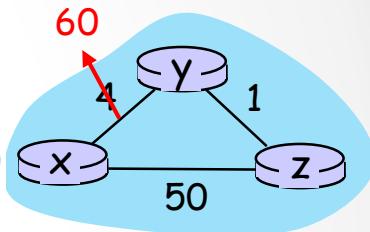
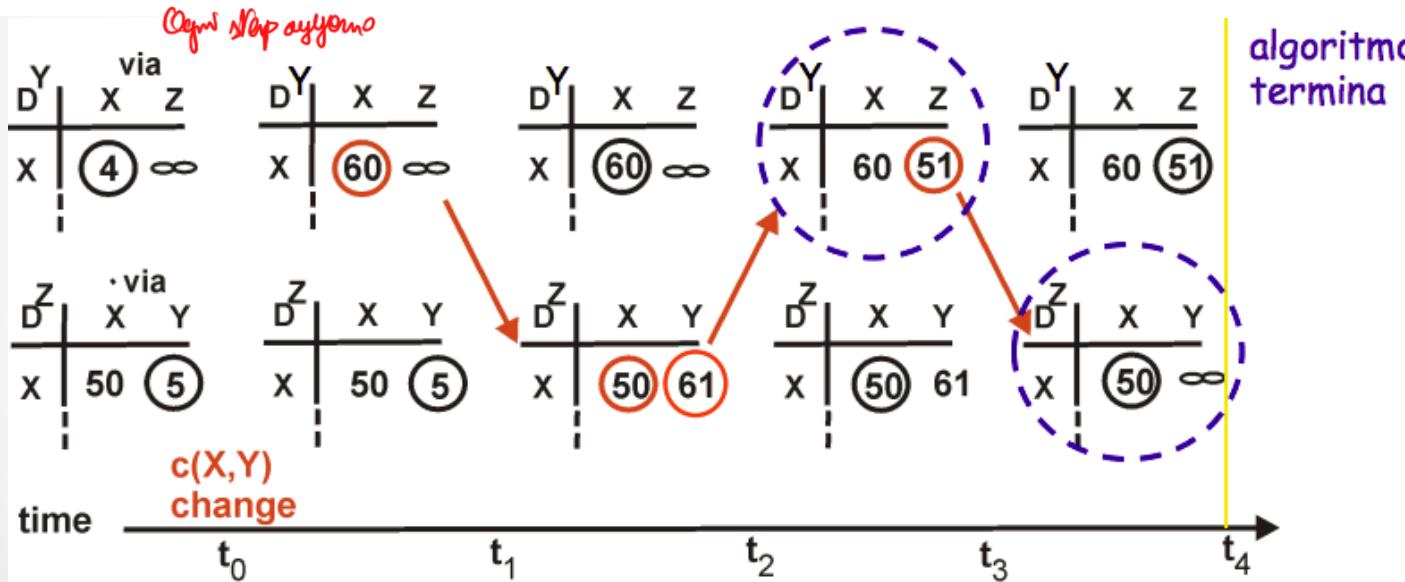


■ Possibile soluzione al problema?

# Distance vector: poisoned reverse

Se z raggiunge x tramite y:

- z dice a y che la sua distanza per x è infinita (così y non andrà a x attraverso z)
- Viene risolto completamente il problema?



- **Problema risolto?**



# Count to Infinity

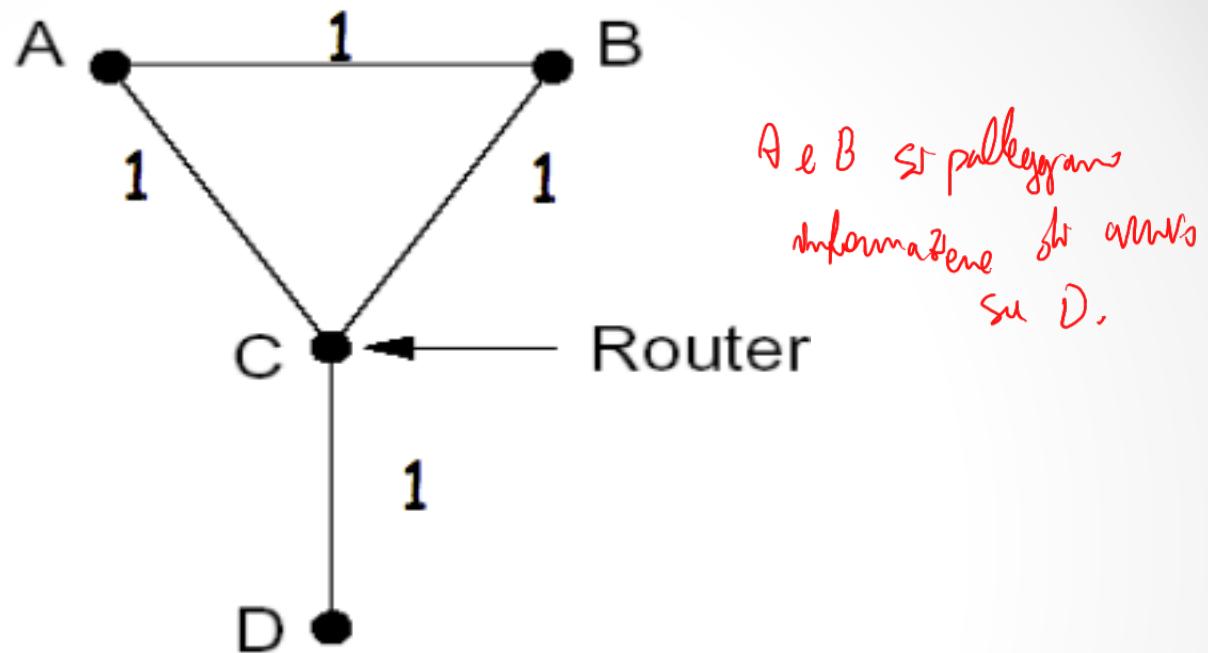
A	B	C	D	E	
X					Siamo tutti che è morto
inf.	2	3	4		B learns A is dead
inf.	inf. $\xrightarrow{2}$	3	4		B reports to C that A's metric is inf.
inf.	inf.	3	4		After 1 exchange
inf.	inf.	inf.	4		After 2 exchanges
inf.	inf.	inf.	inf.		After 3 exchanges



■ **Problema risolto?**



# Fallimento Poisoned Reverse



- Quando il link tra C e D si interrompe, C “setterà” la sua distanza da D ad  $\infty$
- Però, A userà B per andare a D e B userà A per andare a D.
- Dopo questi update, sia A che B riporteranno un nuovo percorso da C a D (diverso da  $\infty$ )

# Comparison of LS and DV algorithms

## message complexity

LS:  $n$  routers,  $O(n^2)$  messages sent

DV: exchange between neighbors;  
convergence time varies

## speed of convergence

LS:  $O(n^2)$  algorithm,  $O(n^2)$  messages

- may have oscillations

DV: convergence time varies

- may have routing loops
- count-to-infinity problem

robustness: what happens if router malfunctions, or is compromised?

LS:

- router can advertise incorrect *link cost*
- each router computes only its *own* table

DV:

- DV router can advertise incorrect *path cost* ("I have a *really* low cost path to everywhere"): black-holing
- each router's table used by others: error propagate thru network